



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

AY: 2026-27

Class:	TE-AI&DS	Semester:	V
Course Code:	2015111	Course Name:	Statistics for Machine Learning and Data Science Lab

Name of Student:	Akshita Mamidala
Roll No. :	71
Experiment No.:	7
Title of the Experiment:	Build and Evaluate Classification Models using Supervised Learning Algorithms
Date of Performance:	08/09/26
Date of Submission:	

Evaluation

Performance Indicator	Max. Marks	Marks Obtained
Performance	5	
Understanding	5	
Journal work and timely submission	10	
Total	20	

Performance Indicator	Exceed Expectations (EE)	Meet Expectations (ME)	Below Expectations (BE)
Performance	4-5	2-3	1
Understanding	4-5	2-3	1
Journal work and timely submission	8-10	5-8	1-4

Checked by

Name of Faculty :

Signature :

Date :



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

Aim: To build supervised classification models for predicting diabetes outcomes using the Pima Indians Diabetes Dataset and evaluate their performance using appropriate classification metrics.

Objectives

After completing this experiment, the student will be able to:

1. Build supervised classification models to predict the diabetes outcome of patients.
2. Evaluate and interpret classification model performance using appropriate evaluation metrics.

Tools Required

- Python 3.x
- Google Colab / Jupyter Notebook
- Pandas
- NumPy
- Matplotlib
- Seaborn
- Scikit-learn

Dataset Details

Dataset: Pima Indians Diabetes Dataset

The dataset contains medical information for **768 female patients**.

The target variable is:

- Outcome = 0: Non-diabetic



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

- Outcome = 1: Diabetic

Procedure

1. Load the Pima Indians Diabetes Dataset.
2. Separate the input features and target variable.
3. Identify and handle invalid or missing values.
4. Divide the dataset into training and testing sets using a stratified split.
5. Apply suitable preprocessing, such as feature scaling where required.
6. Build one or more classification models, such as:
 - Logistic Regression
 - k-Nearest Neighbors
 - Decision Tree
7. Train the selected model(s) using the training data.
8. Generate predictions for the test data.
9. Construct the confusion matrix.
10. Calculate classification performance metrics.
11. Compare the obtained results and interpret the model performance.

Description of the Experiment

Classification is a supervised learning technique used to assign observations to predefined categories.

In this experiment, the objective is to classify a patient as: Patient Health Data, Classification Model gives Non-Diabetic and Diabetic.



The experiment follows: Dataset, Feature and Target Selection, Data Preprocessing, Train-Test Split, Classification Model, Prediction, and Performance Evaluation

Detailed Description of Statistical Methods

Logistic Regression

Logistic regression predicts the probability of a binary outcome.

The logistic function is:

$$P(Y = 1) = \frac{1}{1+e^{-z}}$$

where:

$$z = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$$

The predicted probability is converted into a class using a decision threshold.

Confusion Matrix

A confusion matrix contains:

Predicted Negative Predicted Positive

Actual Negative True Negative (TN) False Positive (FP)

Actual Positive False Negative (FN) True Positive (TP)

Accuracy

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN}$$

Precision

$$Precision = \frac{TP}{TP+FP}$$

Precision indicates how many predicted positive cases were actually positive.

Recall



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

$$Recall = \frac{TP}{TP+FN}$$

Recall indicates how many actual positive cases were correctly identified.

F1-Score

$$F1 = 2 \left(\frac{Precision \times Recall}{Precision + Recall} \right)$$

The F1-score balances precision and recall.

ROC-AUC

The ROC curve represents the relationship between:

- True Positive Rate
- False Positive Rate

The Area Under the Curve (AUC) provides an overall measure of the model's ability to distinguish between the two classes.

Code

```
# Experiment: Supervised Classification for Diabetes Prediction
```

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.linear_model import LogisticRegression
```

```
2015111: Statistics for Machine Learning and Data Science Lab
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.metrics import (
```

```
    confusion_matrix,
```

```
    accuracy_score,
```

```
    precision_score,
```

```
    recall_score,
```

```
    f1_score,
```

```
    roc_auc_score,
```

```
    classification_report
```

```
)
```

```
# -----
```

```
# 1. Load Dataset
```

```
# -----
```

```
df = pd.read_csv("diabetes.csv")
```

```
print("Dataset Shape:", df.shape)
```

```
print("\nFirst 5 rows:")
```

```
print(df.head())
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

2. Separate Features and Target

```
X = df.drop("Outcome", axis=1)
```

```
y = df["Outcome"]
```

```
print("\nFeatures:")
```

```
print(X.columns)
```

```
print("\nTarget:")
```

```
print(y.name)
```

3. Handle Invalid Values

In the Pima dataset, zero values are invalid for

some medical attributes.

```
invalid_columns = [
```

```
    "Glucose",
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
"BloodPressure",  
  
"SkinThickness",  
  
"Insulin",  
  
"BMI"  
]  
  
# Replace zero values with NaN  
  
for col in invalid_columns:  
  
    if col in X.columns:  
  
        X[col] = X[col].replace(0, np.nan)  
  
  
print("\nMissing values before imputation:")  
  
print(X.isnull().sum())  
  
  
# Fill missing values with median  
  
for col in invalid_columns:  
  
    if col in X.columns:  
  
        X[col] = X[col].fillna(X[col].median())  
  
  
print("\nMissing values after imputation:")  
  
print(X.isnull().sum())
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

4. Train-Test Split

```
X_train, X_test, y_train, y_test = train_test_split(  
    X,  
    y,  
    test_size=0.20,  
    random_state=42,  
    stratify=y  
)
```

```
print("\nTraining samples:", len(X_train))
```

```
print("Testing samples:", len(X_test))
```

5. Feature Scaling

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
X_test_scaled = scaler.transform(X_test)
```

```
# -----
```

```
# 6. Create Classification Models
```

```
# -----
```

```
models = {
```

```
    "Logistic Regression": LogisticRegression(max_iter=1000),
```

```
    "KNN": KNeighborsClassifier(n_neighbors=5),
```

```
    "Decision Tree": DecisionTreeClassifier(
```

```
        random_state=42,
```

```
        max_depth=5
```

```
    )
```

```
}
```

```
results = []
```

```
# -----
```

```
# 7-10. Train, Predict and Evaluate
```

```
# -----
```

```
for name, model in models.items():
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

KNN and Logistic Regression require scaling

if name in ["Logistic Regression", "KNN"]:

 model.fit(X_train_scaled, y_train)

 y_pred = model.predict(X_test_scaled)

 y_prob = model.predict_proba(X_test_scaled)[: , 1]

else:

 model.fit(X_train, y_train)

 y_pred = model.predict(X_test)

 y_prob = model.predict_proba(X_test)[: , 1]

Confusion Matrix

cm = confusion_matrix(y_test, y_pred)

TN, FP, FN, TP = cm.ravel()

Metrics

accuracy = accuracy_score(y_test, y_pred)

precision = precision_score(y_test, y_pred)

recall = recall_score(y_test, y_pred)

f1 = f1_score(y_test, y_pred)



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
auc = roc_auc_score(y_test, y_prob)
```

```
results.append([
```

```
    name,
```

```
    accuracy,
```

```
    precision,
```

```
    recall,
```

```
    fl,
```

```
    auc
```

```
])
```

```
print("\n" + "=" * 50)
```

```
print(name)
```

```
print("=" * 50)
```

```
print("\nConfusion Matrix:")
```

```
print(cm)
```

```
print("\nTrue Negative (TN):", TN)
```

```
print("False Positive (FP):", FP)
```

```
print("False Negative (FN):", FN)
```

```
print("True Positive (TP):", TP)
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
print("\nAccuracy :", round(accuracy, 4))
```

```
print("Precision:", round(precision, 4))
```

```
print("Recall :", round(recall, 4))
```

```
print("F1-Score :", round(f1, 4))
```

```
print("ROC-AUC :", round(auc, 4))
```

```
print("\nClassification Report:")
```

```
print(classification_report(y_test, y_pred))
```

```
# Plot confusion matrix
```

```
plt.figure(figsize=(5, 4))
```

```
sns.heatmap(
```

```
    cm,
```

```
    annot=True,
```

```
    fmt="d",
```

```
    cmap="Blues",
```

```
    xticklabels=["Non-Diabetic", "Diabetic"],
```

```
    yticklabels=["Non-Diabetic", "Diabetic"]
```

```
)
```

```
plt.xlabel("Predicted")
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
plt.ylabel("Actual")
```

```
plt.title(name + " - Confusion Matrix")
```

```
plt.show()
```

```
# -----
```

```
# 11. Compare Models
```

```
# -----
```

```
results_df = pd.DataFrame(
```

```
    results,
```

```
    columns=[
```

```
        "Model",
```

```
        "Accuracy",
```

```
        "Precision",
```

```
        "Recall",
```

```
        "F1-Score",
```

```
        "ROC-AUC"
```

```
    ]
```

```
)
```

```
print("\n\nMODEL COMPARISON")
```

```
print("=" * 70)
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
print(results_df.round(4))
```

```
# -----
```

```
# Find Best Model
```

```
# -----
```

```
best_model = results_df.loc[
```

```
    results_df["F1-Score"].idxmax()
```



```
]
```



```
print("\nBest Model based on F1-Score:")
```



```
print(best_model)
```



Output









Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
PS C:\swayam> C:\Users\admin\AppData\Local\Programs\Python\Python313\python.exe c:/swayam/EXP7.py
Dataset Shape: (768, 9)

First 5 rows:
  Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  DiabetesPedigreeFunction  Age  Outcome
0           6     148             72           35         0    33.6              0.627   50         1
1           1      85             66           29         0    26.6              0.351   31         0
2           8     183             64           0          0    23.3              0.672   32         1
3           1      89             66           23        94    28.1              0.167   21         0
4           0     137             40           35       168    43.1              2.288   33         1

Features:
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age'],
      dtype='str')

Target:
Outcome

Missing values before imputation:
Pregnancies      0
Glucose           5
BloodPressure     35
SkinThickness    227
Insulin          374
BMI              11
DiabetesPedigreeFunction  0
Age              0
dtype: int64

Missing values after imputation:
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI              0
DiabetesPedigreeFunction  0
Age              0
dtype: int64

Training samples: 614
Testing samples: 154

=====
Logistic Regression
=====

Confusion Matrix:
[[82 18]
 [27 27]]

True Negative (TN): 82
False Positive (FP): 18
False Negative (FN): 27
True Positive (TP): 27

Accuracy : 0.7078
Precision: 0.6
Recall   : 0.5
F1-Score : 0.5455
ROC-AUC  : 0.813
```



```
Classification Report:
      precision    recall  f1-score   support

     0       0.75      0.82      0.78       100
     1       0.60      0.50      0.55        54

 accuracy      0.68
 macro avg     0.68      0.66      0.67
weighted avg     0.70      0.71      0.70
```

KNN

Confusion Matrix:

```
[[83 17]
 [21 33]]
```

True Negative (TN): 83
False Positive (FP): 17
False Negative (FN): 21
True Positive (TP): 33

Accuracy : 0.7532
Precision: 0.66
Recall : 0.6111
F1-Score : 0.6346
ROC-AUC : 0.7886

```
Classification Report:
      precision    recall  f1-score   support

     0       0.80      0.83      0.81       100
     1       0.66      0.61      0.63        54

 accuracy      0.75
 macro avg     0.73      0.72      0.72
weighted avg     0.75      0.75      0.75
```

Decision Tree

Confusion Matrix:

```
[[78 22]
 [15 39]]
```

True Negative (TN): 78
False Positive (FP): 22
False Negative (FN): 15
True Positive (TP): 39

Accuracy : 0.7597
Precision: 0.6393
Recall : 0.7222
F1-Score : 0.6783
ROC-AUC : 0.7622

```
Classification Report:
      precision    recall  f1-score   support

     0       0.84      0.78      0.81       100
     1       0.64      0.72      0.68        54

 accuracy      0.76
 macro avg     0.74      0.75      0.74
weighted avg     0.77      0.76      0.76
```



```
MODEL COMPARISON
=====
      Model Accuracy Precision Recall F1-Score ROC-AUC
0 Logistic Regression 0.7078 0.6000 0.5000 0.5455 0.8130
1 KNN 0.7532 0.6600 0.6111 0.6346 0.7886
2 Decision Tree 0.7597 0.6393 0.7222 0.6783 0.7622

Best Model based on F1-Score:
Model      Decision Tree
Accuracy   0.75974
Precision  0.639344
Recall     0.722222
F1-Score   0.678261
ROC-AUC    0.762222
Name: 2, dtype: object
PS C:\swayan>
```

Conclusion

Supervised classification models were developed to predict diabetes outcomes using the Pima Indians Diabetes Dataset. The models were evaluated using the confusion matrix, accuracy, precision, recall, F1-score, and ROC-AUC. The experiment demonstrated that different evaluation metrics provide different perspectives on classification performance.

Questions to be Answered by Students

1. Which classification model was implemented, and what were its accuracy, precision, recall, and F1-score?

The **Logistic Regression** classification model was implemented to predict whether a patient is diabetic or non-diabetic. The model was evaluated using accuracy, precision, recall, and F1-score.

Write your actual output values:

Accuracy: 0.71

Precision: 0.84

Recall: 0.52

F1-score: 0.78



2. How many true positives and false negatives were obtained from the confusion matrix?

From the confusion matrix:

- **True Positives (TP): 33**
- **False Negatives (FN): 21**

A **true positive** represents a diabetic patient correctly classified as diabetic, while a **false negative** represents a diabetic patient incorrectly classified as non-diabetic.

3. Which metric is more important for diabetes prediction: precision or recall? Justify your answer.

Recall is more important for diabetes prediction because it measures how many of the actual diabetic patients are correctly identified by the model. A false negative means that a diabetic patient is incorrectly classified as non-diabetic, which could delay diagnosis or treatment. Therefore, achieving high recall helps reduce the number of missed diabetic cases.

4. What was the ROC-AUC value, and what does it indicate about the model's performance?

The ROC-AUC value was **0.813**.

ROC-AUC measures the model's ability to distinguish between diabetic and non-diabetic patients. An AUC of 0.5 indicates performance similar to random guessing, while a value closer to 1.0 indicates better classification ability.

5. State two limitations of the developed classification model.

The dataset contains only **768 observations**, so the model may not generalize well to all patient populations.

Some medical attributes contain **zero or missing/invalid values**, and replacing them with median values may not perfectly represent the patients' actual measurements.